

Лабораторная работа №11 “Реализация криптографических атак с помощью машинного обучения на физически неклонлируемые функции”

Долматович Алина, 858641

Сформулируйте задачу в терминах машинного обучения.

Задача: Обучить модель машинного обучения таким образом, чтобы она смогла предсказывать выходное значение (Response, R) по запросу (Challenge, CH)

```
In [48]: import pandas
import numpy as np
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.tree import DecisionTreeClassifier
import sklearn.cross_validation
from sklearn.metrics import accuracy_score, f1_score, log_loss, recall
import matplotlib.pyplot as pyplot
```

Обучите модель, которая могла бы предсказывать ответы по запросам, которых нет в обучающей выборке.

```
In [2]: def getData(fileName, rowCount=500):
    data = pandas.read_csv(fileName, sep=" ", header=None, nrows=rowCount)
    symbolsCount = int(fileName.replace("Base", "").replace(".txt", ""))

    x = np.array(list(data[0].apply(str).map(lambda x: np.array(map(int, x.split())))))
    y = data[1].values

    trainX, testX, trainY, testY = sklearn.cross_validation.train_test_split(x, y,
                                                                                test_size=0.2)

    return trainX, testX, trainY, testY

trainX, testX, trainY, testY = getData("Base128.txt")
print(trainX.shape, trainY.shape)
print(testX.shape, testY.shape)

((400, 128), (400,))
((100, 128), (100,))
```

Применить как минимум 3 различных алгоритма (например, метод опорных векторов, логистическая регрессия и градиентный бустинг).

```
In [3]: def fitModel(algorithm, x, y):
        return algorithm.fit(x, y)

def getModels(x, y):
    algorithms = [KNeighborsClassifier(), DecisionTreeClassifier(), GradientBoostingClassifier()]
    models = []
    for algorithm in algorithms:
        model = fitModel(algorithm, x, y)
        models.append(model)
    return models

models = getModels(trainX, trainY)
```

```
In [4]: print models

[KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                      metric_params=None, n_jobs=1, n_neighbors=5, p=2,
                      weights='uniform'), DecisionTreeClassifier(class_weight=None,
                      criterion='gini', max_depth=None,
                      max_features=None, max_leaf_nodes=None,
                      min_impurity_decrease=0.0, min_impurity_split=None,
                      min_samples_leaf=1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                      splitter='best'), GradientBoostingClassifier(criterion='friedman_mse',
                      init=None,
                      learning_rate=0.1, loss='deviance', max_depth=3,
                      max_features=None, max_leaf_nodes=None,
                      min_impurity_decrease=0.0, min_impurity_split=None,
                      min_samples_leaf=1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, n_estimators=100,
                      presort='auto', random_state=None, subsample=1.0, verbose=0,
                      warm_start=False)]
```

Какая метрика наиболее подходит для оценки качества алгоритма?

Какой наибольшей доли правильных ответов (Accuracy) удалось достичь?

```
In [5]: def getMetrics(models, testX, testY):
    metrics = [accuracy_score, f1_score, log_loss, recall_score]

    results = dict()

    for model in models:
        print "\n", type(model)
        predict = model.predict(testX)
        metricsValue = []
        for metric in metrics:
            print str(metric), "=====>", metric(testY, predict)
            metricsValue.append(metric(testY, predict))
        results[str(type(model))] = metricsValue
    return results

metrics = getMetrics(models, testX, testY)
```

```
<class 'sklearn.neighbors.classification.KNeighborsClassifier'>
<function accuracy_score at 0x1a1597b5f0> =====> 0.55
<function f1_score at 0x1a1597b8c0> =====> 0.68085106383
<function log_loss at 0x1a1597bc80> =====> 15.5427212408
<function recall_score at 0x1a1597bb18> =====> 0.813559322034

<class 'sklearn.tree.tree.DecisionTreeClassifier'>
<function accuracy_score at 0x1a1597b5f0> =====> 0.58
<function f1_score at 0x1a1597b8c0> =====> 0.671875
<function log_loss at 0x1a1597bc80> =====> 14.5064939812
<function recall_score at 0x1a1597bb18> =====> 0.728813559322

<class 'sklearn.ensemble.gradient_boosting.GradientBoostingClassifier'>
<function accuracy_score at 0x1a1597b5f0> =====> 0.55
<function f1_score at 0x1a1597b8c0> =====> 0.676258992806
<function log_loss at 0x1a1597bc80> =====> 15.5427132449
<function recall_score at 0x1a1597bb18> =====> 0.796610169492
```

Какой размер обучающей выборки необходим, чтобы достигнуть доли правильных ответов минимум 0.95?

```
In [6]: def investigation(rowsCount):
    trainX, testX, trainY, testY = getData("Base128.txt", rowsCount=rowsCount)
    models = getModels(trainX, trainY)
    metrics = getMetrics(models, testX, testY)
    return metrics
```

Как зависит доля правильных ответов от N?

```
In [7]: history = dict()
    nCounts = [100, 1000, 10000, 20000, 30000]
    for n in nCounts:
```

```
print "\n", n
history[n] = investigation(n)
```

100

```
<class 'sklearn.neighbors.classification.KNeighborsClassifier'>
<function accuracy_score at 0x1a1597b5f0> =====> 0.65
<function f1_score at 0x1a1597b8c0> =====> 0.758620689655
<function log_loss at 0x1a1597bc80> =====> 12.0887316577
<function recall_score at 0x1a1597bb18> =====> 0.785714285714

<class 'sklearn.tree.tree.DecisionTreeClassifier'>
<function accuracy_score at 0x1a1597b5f0> =====> 0.7
<function f1_score at 0x1a1597b8c0> =====> 0.8125
<function log_loss at 0x1a1597bc80> =====> 10.3618328178
<function recall_score at 0x1a1597bb18> =====> 0.928571428571

<class 'sklearn.ensemble.gradient_boosting.GradientBoostingClassifier'>
<function accuracy_score at 0x1a1597b5f0> =====> 0.65
<function f1_score at 0x1a1597b8c0> =====> 0.774193548387
<function log_loss at 0x1a1597bc80> =====> 12.0887716376
<function recall_score at 0x1a1597bb18> =====> 0.857142857143
```

1000

```
<class 'sklearn.neighbors.classification.KNeighborsClassifier'>
<function accuracy_score at 0x1a1597b5f0> =====> 0.53
<function f1_score at 0x1a1597b8c0> =====> 0.654411764706
<function log_loss at 0x1a1597bc80> =====> 16.2334887728
<function recall_score at 0x1a1597bb18> =====> 0.760683760684

<class 'sklearn.tree.tree.DecisionTreeClassifier'>
<function accuracy_score at 0x1a1597b5f0> =====> 0.59
<function f1_score at 0x1a1597b8c0> =====> 0.655462184874
<function log_loss at 0x1a1597bc80> =====> 14.1610702354
<function recall_score at 0x1a1597bb18> =====> 0.666666666667

<class 'sklearn.ensemble.gradient_boosting.GradientBoostingClassifier'>
<function accuracy_score at 0x1a1597b5f0> =====> 0.555
<function f1_score at 0x1a1597b8c0> =====> 0.676363636364
<function log_loss at 0x1a1597bc80> =====> 15.3700153649
<function recall_score at 0x1a1597bb18> =====> 0.794871794872
```

10000

```
<class 'sklearn.neighbors.classification.KNeighborsClassifier'>
<function accuracy_score at 0x1a1597b5f0> =====> 0.551
<function f1_score at 0x1a1597b8c0> =====> 0.666666666667
<function log_loss at 0x1a1597bc80> =====> 15.5081240938
<function recall_score at 0x1a1597bb18> =====> 0.711568938193
```

```

<class 'sklearn.tree.tree.DecisionTreeClassifier'>
<function accuracy_score at 0x1a1597b5f0> =====> 0.527
<function f1_score at 0x1a1597b8c0> =====> 0.612929623568
<function log_loss at 0x1a1597bc80> =====> 16.3370143476
<function recall_score at 0x1a1597bb18> =====> 0.593502377179

<class 'sklearn.ensemble.gradient_boosting.GradientBoostingClassifier'>
<function accuracy_score at 0x1a1597b5f0> =====> 0.6275
<function f1_score at 0x1a1597b8c0> =====> 0.756297023225
<function log_loss at 0x1a1597bc80> =====> 12.8659496785
<function recall_score at 0x1a1597bb18> =====> 0.916006339144

```

20000

```

<class 'sklearn.neighbors.classification.KNeighborsClassifier'>
<function accuracy_score at 0x1a1597b5f0> =====> 0.56525
<function f1_score at 0x1a1597b8c0> =====> 0.675982858208
<function log_loss at 0x1a1597bc80> =====> 15.0159441314
<function recall_score at 0x1a1597bb18> =====> 0.726471766119

<class 'sklearn.tree.tree.DecisionTreeClassifier'>
<function accuracy_score at 0x1a1597b5f0> =====> 0.52225
<function f1_score at 0x1a1597b8c0> =====> 0.609122519943
<function log_loss at 0x1a1597bc80> =====> 16.5010809318
<function recall_score at 0x1a1597bb18> =====> 0.596315578694

<class 'sklearn.ensemble.gradient_boosting.GradientBoostingClassifier'>
<function accuracy_score at 0x1a1597b5f0> =====> 0.6245
<function f1_score at 0x1a1597b8c0> =====> 0.757115135834
<function log_loss at 0x1a1597bc80> =====> 12.9695796008
<function recall_score at 0x1a1597bb18> =====> 0.937525030036

```

30000

```

<class 'sklearn.neighbors.classification.KNeighborsClassifier'>
<function accuracy_score at 0x1a1597b5f0> =====> 0.553
<function f1_score at 0x1a1597b8c0> =====> 0.662386706949
<function log_loss at 0x1a1597bc80> =====> 15.4390456082
<function recall_score at 0x1a1597bb18> =====> 0.707638515331

<class 'sklearn.tree.tree.DecisionTreeClassifier'>
<function accuracy_score at 0x1a1597b5f0> =====> 0.5265
<function f1_score at 0x1a1597b8c0> =====> 0.606890826069
<function log_loss at 0x1a1597bc80> =====> 16.3542860014
<function recall_score at 0x1a1597bb18> =====> 0.589833243679

<class 'sklearn.ensemble.gradient_boosting.GradientBoostingClassifier'>
<function accuracy_score at 0x1a1597b5f0> =====> 0.615333333333
<function f1_score at 0x1a1597b8c0> =====> 0.750809760311
<function log_loss at 0x1a1597bc80> =====> 13.2861914479

```

<function recall_score at 0x1a1597bb18> =====> 0.935180204411

Ответы на вопросы представьте в виде графиков.

```
In [47]: originalKeys = history.values()[0].keys()

accuracy_score_values = dict()
f1_score_values = dict()
log_loss_values = dict()
recall_score_values = dict()

for key in originalKeys:
    for n in history:
        values = history[n][key]
        accuracy_score_values[n] = values[0]
        f1_score_values[n] = values[1]
        recall_score_values[n] = values[3]

pyplot.scatter(accuracy_score_values.keys(), accuracy_score_values.values(), c='yellow')
pyplot.scatter(f1_score_values.keys(), f1_score_values.values(), c='green')
pyplot.scatter(recall_score_values.keys(), recall_score_values.values(), c='red')
pyplot.show()
```



