

Лабораторная работа №2. Реализация глубокой нейронной сети

Данные: В работе предлагается использовать набор данных notMNIST, который состоит из изображений размерностью 28×28 первых 10 букв латинского алфавита (A ... J, соответственно). Обучающая выборка содержит порядка 500 тыс. изображений, а тестовая – около 19 тыс.

Данные можно скачать по ссылке:

https://commondatastorage.googleapis.com/books1000/notMNIST_large.tar.gz
(https://commondatastorage.googleapis.com/books1000/notMNIST_large.tar.gz)
(большой набор данных);

https://commondatastorage.googleapis.com/books1000/notMNIST_small.tar.gz
(https://commondatastorage.googleapis.com/books1000/notMNIST_small.tar.gz)
(маленький набор данных);

Описание данных на английском языке доступно по ссылке:

<http://yaroslavvb.blogspot.sg/2011/09/notmnist-dataset.html>
(<http://yaroslavvb.blogspot.sg/2011/09/notmnist-dataset.html>)

```
In [1]: 1 import tensorflow as tf
        2 from tensorflow import keras
        3 import numpy as np
        4 import matplotlib.pyplot as plt
        5
        6 from scipy.io import savemat
        7 import numpy, glob, sys, os
        8 from PIL import Image
        9 from scipy.io import loadmat
```

```
In [2]: 1 def generateDataset(folder, target):
2         max_count = 0
3         for (root, dirs, files) in os.walk(folder):
4             for f in files:
5                 if f.endswith('.png'):
6                     max_count += 1
7             print('Found %s files' % (max_count,))
8             data = numpy.zeros((28,28,max_count))
9             labels = numpy.zeros((max_count,))
10            count = 0
11            for (root, dirs, files) in os.walk(folder):
12                for f in files:
13                    if f.endswith('.png'):
14                        try:
15                            img = Image.open(os.path.join(root,f));
16                            data[:, :, count]=numpy.asarray(img)
17                            surround_folder = os.path.split(root)[-1]
18                            assert len(surround_folder)==1
19                            labels[count]=ord(surround_folder)-ord('A')
20                            count+=1
21                        except:
22                            pass
23
24            print('Saving to ', target)
25            savemat(target,{'images': data[:, :, :count], 'labels': labels
```

```
In [3]: 1 generateDataset("notMNIST_small", "test_dataset.mat")
2         generateDataset("notMNIST_large", "train_dataset.mat")
```

```
Found 18726 files
Saving to test_dataset.mat
Found 529119 files
Saving to train_dataset.mat
```

```
In [4]: 1 def getDataset(file):
2         matfile = loadmat(file)
3         x = matfile['images'] / 255
4         y = matfile['labels']
5         return x.T, y.T
6
7 trainX, trainY = getDataset('train_dataset.mat')
8 print(trainX.shape, trainY.shape)
9
10 testX, testY = getDataset('test_dataset.mat')
11 print(testX.shape, testY.shape)
```

```
(529115, 28, 28) (529115, 1)
(18724, 28, 28) (18724, 1)
```

Задание 1. Реализуйте полносвязную нейронную сеть с помощью библиотеки Tensor Flow. В качестве алгоритма оптимизации можно использовать, например, стохастический градиент (Stochastic Gradient Descent, SGD). Определите количество скрытых слоев от 1 до 5, количество нейронов в каждом из слоев до нескольких сотен, а также их функции активации (кусочно-линейная, сигмоидная, гиперболический тангенс и т.д.).

```
In [5]: 1 model = keras.Sequential([
2         keras.layers.Flatten(input_shape=(28, 28)),
3         keras.layers.Dense(256, activation='relu'),
4         keras.layers.Dense(128, activation='relu'),
5         keras.layers.Dense(64, activation='relu'),
6         keras.layers.Dense(10, activation='softmax')
7     ])
8
9 model.compile(optimizer='adam',
10               loss='sparse_categorical_crossentropy',
11               metrics=['accuracy'])
12
13 model.fit(trainX, trainY, epochs=15)
```

WARNING:tensorflow:From /Users/alinadolmatovich/anaconda2/envs/tensorflow-env/lib/python3.6/site-packages/tensorflow_core/python/ops/resource_variable_ops.py:1630: calling BaseResourceVariable.__init__ (from tensorflow.python.ops.resource_variable_ops) with constraint is deprecated and will be removed in a future version.

Instructions for updating:

If using Keras pass *_constraint arguments to layers.

Train on 529115 samples

Epoch 1/15

529115/529115 [=====] - 199s 375us/sample

- loss: 0.4140 - acc: 0.8734

```
Epoch 2/15
529115/529115 [=====] - 192s 362us/sample
- loss: 0.3290 - acc: 0.8979
Epoch 3/15
529115/529115 [=====] - 190s 360us/sample
- loss: 0.3008 - acc: 0.9061
Epoch 4/15
529115/529115 [=====] - 185s 349us/sample
- loss: 0.2836 - acc: 0.9110
Epoch 5/15
529115/529115 [=====] - 178s 337us/sample
- loss: 0.2704 - acc: 0.9148
Epoch 6/15
529115/529115 [=====] - 179s 338us/sample
- loss: 0.2609 - acc: 0.9177
Epoch 7/15
529115/529115 [=====] - 174s 328us/sample
- loss: 0.2534 - acc: 0.9199
Epoch 8/15
529115/529115 [=====] - 181s 342us/sample
- loss: 0.2470 - acc: 0.9218
Epoch 9/15
529115/529115 [=====] - 168s 318us/sample
- loss: 0.2403 - acc: 0.9236
Epoch 10/15
529115/529115 [=====] - 191s 362us/sample
- loss: 0.2366 - acc: 0.9248
Epoch 11/15
529115/529115 [=====] - 202s 382us/sample
- loss: 0.2324 - acc: 0.9259
Epoch 12/15
529115/529115 [=====] - 179s 337us/sample
- loss: 0.2280 - acc: 0.9273
Epoch 13/15
529115/529115 [=====] - 187s 353us/sample
- loss: 0.2246 - acc: 0.9287
Epoch 14/15
529115/529115 [=====] - 188s 356us/sample
- loss: 0.2213 - acc: 0.9293
Epoch 15/15
529115/529115 [=====] - 181s 342us/sample
- loss: 0.2184 - acc: 0.9304
```

Out[5]: <tensorflow.python.keras.callbacks.History at 0x13face4a8>

Задание 2. Как улучшилась точность классификатора по сравнению с логистической регрессией?

```
In [8]: 1 testLoss, testAcc = model.evaluate(testX, testY, verbose=2)
        2 print('Точность на проверочных данных:', testAcc)
```

18724/18724 - 2s - loss: 0.1378 - acc: 0.9621
Точность на проверочных данных: 0.9621342

Задание 3. Используйте регуляризацию и метод сброса нейронов (dropout) для борьбы с переобучением. Как улучшилось качество классификации?

```
In [9]: model = keras.models.Sequential([
        1 layers.Flatten(input_shape=(28, 28)),
        2 layers.Dense(256, kernel_regularizer=keras.regularizers.l2(0.001), activation='relu'),
        3 layers.Dense(128, kernel_regularizer=keras.regularizers.l2(0.001), activation='relu'),
        4 layers.Dense(64, kernel_regularizer=keras.regularizers.l2(0.001), activation='relu'),
        5 layers.Dense(10, activation='softmax')
        6
        7
        8
        9 model.compile(optimizer='adam',
        10 loss='sparse_categorical_crossentropy',
        11 metrics=['accuracy'])
        12
        13 model.fit(trainX, trainY, epochs=15)
```

Train on 529115 samples

Epoch 1/15

529115/529115 [=====] - 251s 475us/sample
- loss: 0.6194 - acc: 0.8538

Epoch 2/15

529115/529115 [=====] - 264s 499us/sample
- loss: 0.5291 - acc: 0.8673

Epoch 3/15

529115/529115 [=====] - 295s 558us/sample
- loss: 0.5174 - acc: 0.8694

Epoch 4/15

529115/529115 [=====] - 306s 579us/sample
- loss: 0.5108 - acc: 0.8707

Epoch 5/15

529115/529115 [=====] - 291s 550us/sample
- loss: 0.5062 - acc: 0.8709

Epoch 6/15

529115/529115 [=====] - 312s 589us/sample
- loss: 0.5040 - acc: 0.8715

Epoch 7/15

529115/529115 [=====] - 295s 557us/sample
- loss: 0.5038 - acc: 0.8713

Epoch 8/15

529115/529115 [=====] - 283s 536us/sample
- loss: 0.5027 - acc: 0.8714

Epoch 9/15

529115/529115 [=====] - 283s 536us/sample
- loss: 0.5027 - acc: 0.8714

```

529115/529115 [=====] - 244s 460us/sample
- loss: 0.5017 - acc: 0.8715
Epoch 10/15
529115/529115 [=====] - 239s 452us/sample
- loss: 0.4999 - acc: 0.8710
Epoch 11/15
529115/529115 [=====] - 244s 461us/sample
- loss: 0.4994 - acc: 0.8715
Epoch 12/15
529115/529115 [=====] - 239s 451us/sample
- loss: 0.4991 - acc: 0.8710
Epoch 13/15
529115/529115 [=====] - 243s 459us/sample
- loss: 0.4976 - acc: 0.8716
Epoch 14/15
529115/529115 [=====] - 251s 475us/sample
- loss: 0.4980 - acc: 0.8716
Epoch 15/15
529115/529115 [=====] - 249s 470us/sample
- loss: 0.4981 - acc: 0.8712

```

Out[9]: <tensorflow.python.keras.callbacks.History at 0x65c8dde10>

```
In [11]: loss, testRegularizationAcc = regularizationModel.evaluate(testX, te
проверочных данных с регуляризацией:', testRegularizationAcc)
```

```

18724/18724 - 3s - loss: 0.3101 - acc: 0.9301
Точность на проверочных данных с регуляризацией: 0.9300897

```

```
In [12]: 1 dropoutModel = keras.Sequential([
2         keras.layers.Flatten(input_shape=(28, 28)),
3         keras.layers.Dropout(0.5),
4         keras.layers.Dense(256, activation='relu'),
5         keras.layers.Dropout(0.5),
6         keras.layers.Dense(128, activation='relu'),
7         keras.layers.Dropout(0.5),
8         keras.layers.Dense(64, activation='relu'),
9         keras.layers.Dropout(0.5),
10        keras.layers.Dense(10, activation='softmax')
11    ])
12
13 dropoutModel.compile(optimizer='adam',
14                      loss='sparse_categorical_crossentropy',
15                      metrics=['accuracy'])
16
17 dropoutModel.fit(trainX, trainY, epochs=15)
```

```

Train on 529115 samples
Epoch 1/15

```

```

529115/529115 [=====] - 284s 536us/sample
- loss: 0.7617 - acc: 0.7801
Epoch 2/15
529115/529115 [=====] - 284s 537us/sample
- loss: 0.6313 - acc: 0.8222
Epoch 3/15
529115/529115 [=====] - 280s 529us/sample
- loss: 0.6077 - acc: 0.8285
Epoch 4/15
529115/529115 [=====] - 280s 529us/sample
- loss: 0.5949 - acc: 0.8326
Epoch 5/15
529115/529115 [=====] - 287s 542us/sample
- loss: 0.5863 - acc: 0.8340
Epoch 6/15
529115/529115 [=====] - 279s 527us/sample
- loss: 0.5816 - acc: 0.8358
Epoch 7/15
529115/529115 [=====] - 306s 579us/sample
- loss: 0.5772 - acc: 0.8367
Epoch 8/15
529115/529115 [=====] - 284s 537us/sample
- loss: 0.5737 - acc: 0.8382
Epoch 9/15
529115/529115 [=====] - 314s 593us/sample
- loss: 0.5706 - acc: 0.8389
Epoch 10/15
529115/529115 [=====] - 365s 691us/sample
- loss: 0.5695 - acc: 0.8391
Epoch 11/15
529115/529115 [=====] - 381s 720us/sample
- loss: 0.5662 - acc: 0.8403
Epoch 12/15
529115/529115 [=====] - 435s 822us/sample
- loss: 0.5639 - acc: 0.8414- lo
Epoch 13/15
529115/529115 [=====] - 406s 767us/sample
- loss: 0.5646 - acc: 0.8411
Epoch 14/15
529115/529115 [=====] - 325s 615us/sample
- loss: 0.5609 - acc: 0.8420
Epoch 15/15
529115/529115 [=====] - 344s 650us/sample
- loss: 0.5606 - acc: 0.8418

```

Out[12]: <tensorflow.python.keras.callbacks.History at 0x65c884fd0>

```

In [13]: 1 testDropoutLoss, testDropoutAcc = dropoutModel.evaluate(testX,
          2 print('Точность на проверочных данных с dropout\''ом:', testDrop

```

```

18724/18724 - 4s - loss: 0.2280 - acc: 0.9383
Точность на проверочных данных с dropout\''ом: 0.93826103

```

Задание 4. Воспользуйтесь динамически изменяемой скоростью обучения (learning rate). Наилучшая точность, достигнутая с помощью данной модели составляет 97.1%. Какую точность демонстрирует Ваша реализованная модель?

```
In [14]: 1 model = keras.Sequential([
2         keras.layers.Flatten(input_shape=(28, 28)),
3         keras.layers.Dense(256, activation='relu'),
4         keras.layers.Dense(128, activation='relu'),
5         keras.layers.Dense(64, activation='relu'),
6         keras.layers.Dense(10, activation='softmax')
7     ])
8
9     adamOptimizer = keras.optimizers.Adam(learning_rate=0.01)
10    model.compile(optimizer=adamOptimizer,
11                  loss='sparse_categorical_crossentropy',
12                  metrics=['accuracy'])
13
14    model.fit(trainX, trainY, epochs=15)
```

Train on 529115 samples

Epoch 1/15

529115/529115 [=====] - 238s 450us/sample
- loss: 0.5622 - acc: 0.8365

Epoch 2/15

529115/529115 [=====] - 227s 429us/sample
- loss: 0.5079 - acc: 0.8517

Epoch 3/15

529115/529115 [=====] - 211s 398us/sample
- loss: 0.4919 - acc: 0.8560

Epoch 4/15

529115/529115 [=====] - 216s 408us/sample
- loss: 0.4838 - acc: 0.8582

Epoch 5/15

529115/529115 [=====] - 226s 426us/sample
- loss: 0.4795 - acc: 0.8593

Epoch 6/15

529115/529115 [=====] - 225s 425us/sample
- loss: 0.4743 - acc: 0.8611

Epoch 7/15

529115/529115 [=====] - 268s 506us/sample
- loss: 0.4696 - acc: 0.8624

Epoch 8/15

529115/529115 [=====] - 246s 464us/sample
- loss: 0.4654 - acc: 0.8628

Epoch 9/15

529115/529115 [=====] - 221s 418us/sample
- loss: 0.4641 - acc: 0.8621

Epoch 10/15

529115/529115 [=====] - 249s 471us/sample


```
- loss: 0.4596 - acc: 0.8641
Epoch 11/15
529115/529115 [=====] - 254s 480us/sample
- loss: 0.4609 - acc: 0.8635
Epoch 12/15
529115/529115 [=====] - 243s 459us/sample
- loss: 0.4555 - acc: 0.8657
Epoch 13/15
529115/529115 [=====] - 238s 449us/sample
- loss: 0.4509 - acc: 0.8673
Epoch 14/15
529115/529115 [=====] - 281s 531us/sample
- loss: 0.4529 - acc: 0.8657
Epoch 15/15
529115/529115 [=====] - 250s 473us/sample
- loss: 0.4518 - acc: 0.8658
```

Out[14]: <tensorflow.python.keras.callbacks.History at 0x652009e10>

```
In [15]: 1 testDropoutLoss, testDropoutAcc = model.evaluate(testX, testY,
          2 print('Точность на проверочных данных с learning rate оптимизац
```

```
18724/18724 - 4s - loss: 0.2816 - acc: 0.9214
Точность на проверочных данных с learning rate оптимизацией: 0.921
38433
```