

Лабораторная работа №3. Реализация сверточной нейронной сети

Данные: В работе предлагается использовать набор данных notMNIST, который состоит из изображений размерностью 28×28 первых 10 букв латинского алфавита (A ... J, соответственно). Обучающая выборка содержит порядка 500 тыс. изображений, а тестовая – около 19 тыс.

Данные можно скачать по ссылке:

https://commondatastorage.googleapis.com/books1000/notMNIST_large.tar.gz

(https://commondatastorage.googleapis.com/books1000/notMNIST_large.tar.gz)

(большой набор данных);

https://commondatastorage.googleapis.com/books1000/notMNIST_small.tar.gz

(https://commondatastorage.googleapis.com/books1000/notMNIST_small.tar.gz)

(маленький набор данных); Описание данных на английском языке доступно по

ссылке: <http://yaroslavvb.blogspot.sg/2011/09/notmnist-dataset.html>

(<http://yaroslavvb.blogspot.sg/2011/09/notmnist-dataset.html>)

```
In [0]: 1 import tensorflow as tf
        2 from scipy.io import loadmat
        3 import numpy as np
        4 from keras.utils import to_categorical
        5 from tensorflow.keras import layers, models, Sequential
```

Задание 1. Реализуйте нейронную сеть с двумя сверточными слоями, и одним полносвязным с нейронами с кусочно-линейной функцией активации. Какова точность построенной модели?

```
In [0]: 1 def getDataset(file):
        2     matfile = loadmat(file)
        3     x = matfile['images'] / 255
        4     y = matfile['labels']
        5     x, y = x.T, y.T
        6     x = x.reshape(x.shape[0], x.shape[1], x.shape[2], 1)
        7     y = to_categorical(y)
        8     return x, y
```

```
In [11]: 1 trainX, trainY = getDataset("/content/drive/My Drive/Collab Data/
2         print(trainX.shape, trainY.shape)
3         testX, testY = getDataset("/content/drive/My Drive/Collab Data/
4         print(testX.shape, testY.shape)
```

```
(529115, 28, 28, 1) (529115, 1)
(18724, 28, 28, 1) (18724, 1)
```

```
In [14]: 1 model = Sequential([
2         layers.Conv2D(64, kernel_size=3, activation='relu', input_shape=(28, 28, 1)),
3         layers.Conv2D(32, kernel_size=3, activation='relu'),
4         layers.Flatten(),
5         layers.Dense(10, activation='softmax')
6     ])
7
8 model.compile(optimizer='adam', loss='sparse_categorical_crossentropy')
9
10 model.fit(trainX, trainY, validation_data=(testX, testY), epochs=15)
```

Train on 529115 samples, validate on 18724 samples

Epoch 1/15

529115/529115 [=====] - 52s 98us/sample -

loss: 0.3813 - acc: 0.8950 - val_loss: 0.1521 - val_acc: 0.9593

Epoch 2/15

529115/529115 [=====] - 52s 99us/sample -

loss: 0.3055 - acc: 0.9151 - val_loss: 0.1340 - val_acc: 0.9636

Epoch 3/15

529115/529115 [=====] - 52s 99us/sample -

loss: 0.2730 - acc: 0.9237 - val_loss: 0.1514 - val_acc: 0.9601

Epoch 4/15

529115/529115 [=====] - 52s 98us/sample -

loss: 0.2502 - acc: 0.9300 - val_loss: 0.1371 - val_acc: 0.9629

Epoch 5/15

529115/529115 [=====] - 52s 98us/sample -

loss: 0.2318 - acc: 0.9345 - val_loss: 0.1356 - val_acc: 0.9632

Epoch 6/15

529115/529115 [=====] - 52s 98us/sample -

loss: 0.2173 - acc: 0.9382 - val_loss: 0.1403 - val_acc: 0.9622

Epoch 7/15

529115/529115 [=====] - 52s 99us/sample -

loss: 0.2056 - acc: 0.9412 - val_loss: 0.1465 - val_acc: 0.9623

Epoch 8/15

529115/529115 [=====] - 52s 99us/sample -

loss: 0.1951 - acc: 0.9440 - val_loss: 0.1500 - val_acc: 0.9614

Epoch 9/15

529115/529115 [=====] - 52s 98us/sample -

loss: 0.1859 - acc: 0.9460 - val_loss: 0.1556 - val_acc: 0.9598

Epoch 10/15

529115/529115 [=====] - 52s 98us/sample -

loss: 0.1784 - acc: 0.9480 - val_loss: 0.1649 - val_acc: 0.9589

Epoch 11/15

```

Epoch 11/15
529115/529115 [=====] - 52s 98us/sample -
loss: 0.1716 - acc: 0.9501 - val_loss: 0.1748 - val_acc: 0.9579
Epoch 12/15
529115/529115 [=====] - 52s 98us/sample -
loss: 0.1652 - acc: 0.9516 - val_loss: 0.1793 - val_acc: 0.9568
Epoch 13/15
529115/529115 [=====] - 52s 98us/sample -
loss: 0.1597 - acc: 0.9530 - val_loss: 0.1783 - val_acc: 0.9592
Epoch 14/15
529115/529115 [=====] - 52s 99us/sample -
loss: 0.1548 - acc: 0.9545 - val_loss: 0.1904 - val_acc: 0.9559
Epoch 15/15
529115/529115 [=====] - 52s 98us/sample -
loss: 0.1498 - acc: 0.9558 - val_loss: 0.1948 - val_acc: 0.9586

```

Out [14]: <tensorflow.python.keras.callbacks.History at 0x7fddb1af5c0>

```

In [16]: 1 testLoss, testAcc = model.evaluate(testX, testY, verbose=2)
        2 print('Точность на проверочных данных:', testAcc)

```

```

18724/18724 - 1s - loss: 0.1948 - acc: 0.9586
Точность на проверочных данных: 0.9586093

```

Задание 2. Замените один из сверточных слоев на слой, реализующий операцию пулинга (Pooling) с функцией максимума или среднего. Как это повлияло на точность классификатора?

```

In [17]: 1 modelWithPooling = Sequential([
        2     layers.Conv2D(64, kernel_size=3, activation='relu', input_s
        3     layers.MaxPooling2D((2, 2)),
        4     layers.Flatten(),
        5     layers.Dense(10, activation='softmax')
        6 ])
        7
        8 modelWithPooling.compile(optimizer='adam', loss='sparse_categor
        9
        10 modelWithPooling.fit(trainX, trainY, validation_data=(testX, te

```

Train on 529115 samples, validate on 18724 samples

```

Epoch 1/15
529115/529115 [=====] - 50s 94us/sample -
loss: 0.4351 - acc: 0.8830 - val_loss: 0.1937 - val_acc: 0.9501
Epoch 2/15
529115/529115 [=====] - 47s 89us/sample -
loss: 0.3723 - acc: 0.8998 - val_loss: 0.1731 - val_acc: 0.9541
Epoch 3/15
529115/529115 [=====] - 48s 90us/sample -
loss: 0.3494 - acc: 0.9053 - val_loss: 0.1665 - val_acc: 0.9560
Epoch 4/15

```

```

529115/529115 [=====] - 48s 90us/sample -
loss: 0.3350 - acc: 0.9088 - val_loss: 0.1646 - val_acc: 0.9566
Epoch 5/15
529115/529115 [=====] - 47s 89us/sample -
loss: 0.3243 - acc: 0.9116 - val_loss: 0.1615 - val_acc: 0.9584
Epoch 6/15
529115/529115 [=====] - 47s 90us/sample -
loss: 0.3168 - acc: 0.9132 - val_loss: 0.1600 - val_acc: 0.9581
Epoch 7/15
529115/529115 [=====] - 47s 89us/sample -
loss: 0.3101 - acc: 0.9149 - val_loss: 0.1565 - val_acc: 0.9586
Epoch 8/15
529115/529115 [=====] - 47s 88us/sample -
loss: 0.3053 - acc: 0.9159 - val_loss: 0.1600 - val_acc: 0.9576
Epoch 9/15
529115/529115 [=====] - 47s 89us/sample -
loss: 0.3009 - acc: 0.9171 - val_loss: 0.1610 - val_acc: 0.9579
Epoch 10/15
529115/529115 [=====] - 48s 91us/sample -
loss: 0.2972 - acc: 0.9181 - val_loss: 0.1594 - val_acc: 0.9568
Epoch 11/15
529115/529115 [=====] - 47s 90us/sample -
loss: 0.2938 - acc: 0.9187 - val_loss: 0.1602 - val_acc: 0.9564
Epoch 12/15
529115/529115 [=====] - 47s 89us/sample -
loss: 0.2913 - acc: 0.9191 - val_loss: 0.1607 - val_acc: 0.9573
Epoch 13/15
529115/529115 [=====] - 47s 89us/sample -
loss: 0.2890 - acc: 0.9199 - val_loss: 0.1639 - val_acc: 0.9557
Epoch 14/15
529115/529115 [=====] - 48s 90us/sample -
loss: 0.2867 - acc: 0.9205 - val_loss: 0.1625 - val_acc: 0.9575
Epoch 15/15
529115/529115 [=====] - 47s 89us/sample -
loss: 0.2848 - acc: 0.9208 - val_loss: 0.1619 - val_acc: 0.9567

```

Out[17]: <tensorflow.python.keras.callbacks.History at 0x7fddb3eeeb38>

```

In [18]: 1 testWithPoolingLoss, testWithPoolingAcc = modelWithPooling.eval
         2 print('Точность на проверочных данных с пулингом:', testWithPool

```

```

18724/18724 - 1s - loss: 0.1619 - acc: 0.9567
Точность на проверочных данных с пулингом: 0.95674

```

Задание 3. Реализуйте классическую архитектуру сверточных сетей LeNet-5 (<http://yann.lecun.com/exdb/lenet/> (<http://yann.lecun.com/exdb/lenet/>)).

```

In [19]: 1 leNet5model = Sequential([
         2     layers.Conv2D(filters=6, kernel_size=3, activation='relu',
         3     layers.AveragePooling2D(),
         4     layers.Conv2D(filters=16, kernel_size=3, activation='relu')
         5     layers.AveragePooling2D(),

```

```

6     layers.Flatten(),
7     layers.Dense(units=120, activation='relu'),
8     layers.Dense(units=84, activation='relu'),
9     layers.Dense(units=10, activation = 'softmax')
10 ]))
11
12 leNet5model.compile(optimizer='adam', loss='sparse_categorical_
13
14 leNet5model.fit(trainX, trainY, validation_data=(testX, testY),

```

Train on 529115 samples, validate on 18724 samples

Epoch 1/15

529115/529115 [=====] - 49s 93us/sample -
loss: 0.4170 - acc: 0.8726 - val_loss: 0.1603 - val_acc: 0.9530

Epoch 2/15

529115/529115 [=====] - 48s 91us/sample -
loss: 0.3040 - acc: 0.9059 - val_loss: 0.1341 - val_acc: 0.9592

Epoch 3/15

529115/529115 [=====] - 48s 91us/sample -
loss: 0.2758 - acc: 0.9140 - val_loss: 0.1201 - val_acc: 0.9630

Epoch 4/15

529115/529115 [=====] - 48s 90us/sample -
loss: 0.2587 - acc: 0.9191 - val_loss: 0.1198 - val_acc: 0.9651

Epoch 5/15

529115/529115 [=====] - 48s 90us/sample -
loss: 0.2476 - acc: 0.9223 - val_loss: 0.1181 - val_acc: 0.9647

Epoch 6/15

529115/529115 [=====] - 48s 91us/sample -
loss: 0.2382 - acc: 0.9249 - val_loss: 0.1159 - val_acc: 0.9656

Epoch 7/15

529115/529115 [=====] - 49s 92us/sample -
loss: 0.2315 - acc: 0.9267 - val_loss: 0.1096 - val_acc: 0.9677

Epoch 8/15

529115/529115 [=====] - 48s 91us/sample -
loss: 0.2251 - acc: 0.9288 - val_loss: 0.1109 - val_acc: 0.9672

Epoch 9/15

529115/529115 [=====] - 48s 90us/sample -
loss: 0.2200 - acc: 0.9301 - val_loss: 0.1083 - val_acc: 0.9682

Epoch 10/15

529115/529115 [=====] - 48s 90us/sample -
loss: 0.2158 - acc: 0.9314 - val_loss: 0.1093 - val_acc: 0.9678

Epoch 11/15

529115/529115 [=====] - 48s 90us/sample -
loss: 0.2113 - acc: 0.9327 - val_loss: 0.1108 - val_acc: 0.9688

Epoch 12/15

529115/529115 [=====] - 48s 90us/sample -
loss: 0.2071 - acc: 0.9338 - val_loss: 0.1110 - val_acc: 0.9675

Epoch 13/15

529115/529115 [=====] - 48s 90us/sample -
loss: 0.2041 - acc: 0.9348 - val_loss: 0.1129 - val_acc: 0.9674

Epoch 14/15

```
529115/529115 [=====] - 49s 92us/sample -  
loss: 0.2010 - acc: 0.9355 - val_loss: 0.1109 - val_acc: 0.9682  
Epoch 15/15  
529115/529115 [=====] - 48s 90us/sample -  
loss: 0.1984 - acc: 0.9363 - val_loss: 0.1137 - val_acc: 0.9668
```

Out[19]: <tensorflow.python.keras.callbacks.History at 0x7fddb3f58860>

```
In [20]: 1 testLeNetLoss, testLeNetAcc = leNet5model.evaluate(testX, testY  
2         print('Точность на проверочных данных с пулингом:', testLeNetAcc
```

```
18724/18724 - 1s - loss: 0.1137 - acc: 0.9668  
Точность на проверочных данных с пулингом: 0.966834
```